

```
In [95]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score
import seaborn as sns
import numpy as np
```

```
In [96]: glass_data = pd.read_csv('glass.csv')
glass_data
```

```
Out[96]:
```

| | RI | Na | Mg | Al | Si | K | Ca | Ba | Fe | Type |
|-----|---------|-------|------|------|-------|------|------|------|-----|------|
| 0 | 1.52101 | 13.64 | 4.49 | 1.10 | 71.78 | 0.06 | 8.75 | 0.00 | 0.0 | 1 |
| 1 | 1.51761 | 13.89 | 3.60 | 1.36 | 72.73 | 0.48 | 7.83 | 0.00 | 0.0 | 1 |
| 2 | 1.51618 | 13.53 | 3.55 | 1.54 | 72.99 | 0.39 | 7.78 | 0.00 | 0.0 | 1 |
| 3 | 1.51766 | 13.21 | 3.69 | 1.29 | 72.61 | 0.57 | 8.22 | 0.00 | 0.0 | 1 |
| 4 | 1.51742 | 13.27 | 3.62 | 1.24 | 73.08 | 0.55 | 8.07 | 0.00 | 0.0 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 209 | 1.51623 | 14.14 | 0.00 | 2.88 | 72.61 | 0.08 | 9.18 | 1.06 | 0.0 | 7 |
| 210 | 1.51685 | 14.92 | 0.00 | 1.99 | 73.06 | 0.00 | 8.40 | 1.59 | 0.0 | 7 |
| 211 | 1.52065 | 14.36 | 0.00 | 2.02 | 73.42 | 0.00 | 8.44 | 1.64 | 0.0 | 7 |
| 212 | 1.51651 | 14.38 | 0.00 | 1.94 | 73.61 | 0.00 | 8.48 | 1.57 | 0.0 | 7 |
| 213 | 1.51711 | 14.23 | 0.00 | 2.08 | 73.36 | 0.00 | 8.62 | 1.67 | 0.0 | 7 |

214 rows × 10 columns

```
In [97]: # initial analysis
glass_data.shape
```

```
Out[97]: (214, 10)
```

```
In [98]: glass_data.isna().sum()
```

```
Out[98]: RI      0
Na      0
Mg      0
Al      0
Si      0
K      0
Ca      0
Ba      0
Fe      0
Type    0
dtype: int64
```

```
In [99]: glass_data.Type.value_counts()
```

```
Out[99]: 2    76  
         1    70  
         7    29  
         3    17  
         5    13  
         6     9  
         Name: Type, dtype: int64
```

```
In [100]: glass_data.dtypes
```

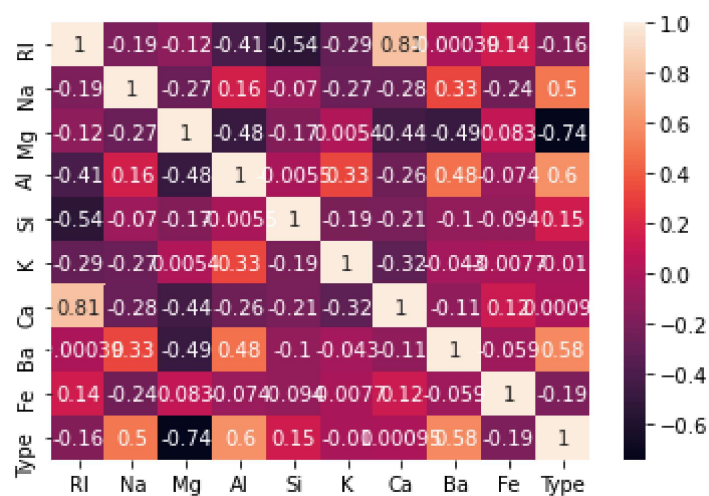
```
Out[100]: RI      float64  
          Na      float64  
          Mg      float64  
          Al      float64  
          Si      float64  
          K       float64  
          Ca      float64  
          Ba      float64  
          Fe      float64  
          Type    int64  
          dtype: object
```

```
In [104]: round(glass_data.corr(),2)
```

```
Out[104]:
```

| | RI | Na | Mg | Al | Si | K | Ca | Ba | Fe | Type |
|------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| RI | 1.00 | -0.19 | -0.12 | -0.41 | -0.54 | -0.29 | 0.81 | -0.00 | 0.14 | -0.16 |
| Na | -0.19 | 1.00 | -0.27 | 0.16 | -0.07 | -0.27 | -0.28 | 0.33 | -0.24 | 0.50 |
| Mg | -0.12 | -0.27 | 1.00 | -0.48 | -0.17 | 0.01 | -0.44 | -0.49 | 0.08 | -0.74 |
| Al | -0.41 | 0.16 | -0.48 | 1.00 | -0.01 | 0.33 | -0.26 | 0.48 | -0.07 | 0.60 |
| Si | -0.54 | -0.07 | -0.17 | -0.01 | 1.00 | -0.19 | -0.21 | -0.10 | -0.09 | 0.15 |
| K | -0.29 | -0.27 | 0.01 | 0.33 | -0.19 | 1.00 | -0.32 | -0.04 | -0.01 | -0.01 |
| Ca | 0.81 | -0.28 | -0.44 | -0.26 | -0.21 | -0.32 | 1.00 | -0.11 | 0.12 | 0.00 |
| Ba | -0.00 | 0.33 | -0.49 | 0.48 | -0.10 | -0.04 | -0.11 | 1.00 | -0.06 | 0.58 |
| Fe | 0.14 | -0.24 | 0.08 | -0.07 | -0.09 | -0.01 | 0.12 | -0.06 | 1.00 | -0.19 |
| Type | -0.16 | 0.50 | -0.74 | 0.60 | 0.15 | -0.01 | 0.00 | 0.58 | -0.19 | 1.00 |

```
In [105]: # Data visualization
sns.heatmap(glass_data.corr(), annot=True)
plt.show()
```



MODEL BUILDING

```
In [106]: X = glass_data.drop(['Type'], axis=1)
y = glass_data['Type']
```

In [107]:

X

Out[107]:

| | RI | Na | Mg | Al | Si | K | Ca | Ba | Fe |
|-----|---------|-------|------|------|-------|------|------|------|-----|
| 0 | 1.52101 | 13.64 | 4.49 | 1.10 | 71.78 | 0.06 | 8.75 | 0.00 | 0.0 |
| 1 | 1.51761 | 13.89 | 3.60 | 1.36 | 72.73 | 0.48 | 7.83 | 0.00 | 0.0 |
| 2 | 1.51618 | 13.53 | 3.55 | 1.54 | 72.99 | 0.39 | 7.78 | 0.00 | 0.0 |
| 3 | 1.51766 | 13.21 | 3.69 | 1.29 | 72.61 | 0.57 | 8.22 | 0.00 | 0.0 |
| 4 | 1.51742 | 13.27 | 3.62 | 1.24 | 73.08 | 0.55 | 8.07 | 0.00 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 209 | 1.51623 | 14.14 | 0.00 | 2.88 | 72.61 | 0.08 | 9.18 | 1.06 | 0.0 |
| 210 | 1.51685 | 14.92 | 0.00 | 1.99 | 73.06 | 0.00 | 8.40 | 1.59 | 0.0 |
| 211 | 1.52065 | 14.36 | 0.00 | 2.02 | 73.42 | 0.00 | 8.44 | 1.64 | 0.0 |
| 212 | 1.51651 | 14.38 | 0.00 | 1.94 | 73.61 | 0.00 | 8.48 | 1.57 | 0.0 |
| 213 | 1.51711 | 14.23 | 0.00 | 2.08 | 73.36 | 0.00 | 8.62 | 1.67 | 0.0 |

214 rows × 9 columns

In [108]:

y

Out[108]:

```

0      1
1      1
2      1
3      1
4      1
..
209    7
210    7
211    7
212    7
213    7

```

Name: Type, Length: 214, dtype: int64

In [109]:

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.20, random_state

In [110]:

X_train.shape , y_train.shape

Out[110]:

((171, 9), (171,))

In [111]:

X_test.shape, y_test.shape

Out[111]:

((43, 9), (43,))

Model training | Testing | Evaluation without NORMALIZATION

Generating a Model with K=3

```
In [112]: knn_model = KNeighborsClassifier(n_neighbors=3)
knn_model.fit(X_train,y_train)
```

```
Out[112]: KNeighborsClassifier(n_neighbors=3)
```

```
In [116]: y_test_pred = knn_model.predict(X_test)
y_test_pred
print("Accuracy score : " , round(accuracy_score(y_test,y_test_pred),4))
```

```
Accuracy score : 0.7907
```

Generating a Model with K=5

```
In [117]: knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train,y_train)
```

```
Out[117]: KNeighborsClassifier()
```

```
In [118]: y_test_pred = knn_model.predict(X_test)
y_test_pred
print("Accuracy score : " , round(accuracy_score(y_test,y_test_pred),4))
```

```
Accuracy score : 0.7907
```

Model Training | Testing | Evaluation with NORMALIZATION

Generating a Model with K = 3

```
In [119]: scaler = StandardScaler()
scaled_X = scaler.fit_transform(X)
```

```
In [120]: X_train,X_test,y_train,y_test = train_test_split(scaled_X,y,test_size=0.20, random_state=42)
X_train.shape , y_train.shape,X_test.shape,y_test.shape
```

```
Out[120]: ((171, 9), (171,), (43, 9), (43,))
```

In [121]: X_train

```
Out[121]: array([[ -2.38151559,  4.87563749, -1.86551055, ..., -1.62482241,
        -0.35287683, -0.5864509 ],
       [ -0.48697606, -0.90567874,  0.13583231, ..., -0.29367195,
        -0.35287683, -0.5864509 ],
       [ -0.15691691, -0.26740179,  0.84464125, ..., -0.36410319,
        -0.35287683,  2.29388828],
       ...,
       [ -0.08430389, -0.095558 ,  0.80989571, ..., -0.58948316,
        -0.35287683, -0.5864509 ],
       [ -0.62890149, -0.45152014,  0.49718589, ..., -0.26549945,
        -0.35287683, -0.5864509 ],
       [ -0.81373462, -0.47606926,  0.62226982, ..., -0.70217315,
        -0.35287683, -0.5864509 ]])
```

In [122]: X_test

```
Out[122]: array([[ 2.56607109e+00,  3.58600602e-01, -1.86551055e+00,
        -1.70460232e-01, -1.82589947e+00, -4.71910254e-01,
         3.15745888e+00, -3.52876828e-01,  4.42241664e-01],
       [ -8.83047040e-01, -2.06029009e-01,  5.59727851e-01,
         5.03781754e-02,  7.75254394e-01, -1.79901870e-01,
        -6.52871279e-01, -3.52876828e-01, -5.86450902e-01],
       [  3.31570637e-01,  4.69071613e-01, -1.90775722e-01,
        -5.11755954e-01,  1.41142259e-01, -7.63918639e-01,
         5.72632323e-01, -3.52876828e-01, -5.86450902e-01],
       [ -2.95541751e-01, -6.97011279e-01,  5.66676959e-01,
        -6.12137048e-01,  8.01136522e-01,  9.67376522e-02,
        -2.86628823e-01, -3.52876828e-01, -5.86450902e-01],
       [ -7.04815098e-01,  8.98681099e-01, -1.86551055e+00,
         2.88112504e+00, -5.29737012e-02, -6.40967740e-01,
         1.57087998e-01,  1.78397794e+00, -5.86450902e-01],
       [  4.37189566e-01,  1.41421248e+00, -1.86551055e+00,
        -1.77655774e+00,  1.07289887e+00, -7.63918639e-01,
         1.59388532e+00, -3.52876828e-01, -5.86450902e-01],
       [ -4.70473101e-01, -6.72462165e-01,  6.22269816e-01,
         3.70841337e-01,  5.30777013e-02,  1.72581064e-01]
```

In [123]: knn_model = KNeighborsClassifier(n_neighbors=3)
knn_model.fit(X_train,y_train)

Out[123]: KNeighborsClassifier(n_neighbors=3)

In [124]: y_test_pred = knn_model.predict(X_test)
y_test_pred
print("Accuracy score : " , round(accuracy_score(y_test,y_test_pred),4))

Accuracy score : 0.7209

Generating a Model with K = 5

```
In [125]: knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train,y_train)
```

```
Out[125]: KNeighborsClassifier()
```

```
In [126]: y_test_pred = knn_model.predict(X_test)
y_test_pred
print("Accuracy score : " , round(accuracy_score(y_test,y_test_pred),4))
```

```
Accuracy score : 0.7442
```

How to pickup Optimum no. of K?

```
In [127]: import warnings
warnings.filterwarnings('ignore')
```

```
In [137]: neighbors = list(range(1,25))
cv_scores = []

for i in neighbors:
    knn_model = KNeighborsClassifier(n_neighbors = i)
    cv_score = cross_val_score(estimator = knn_model, X = scaled_X, y=y, cv = 5)
    cv_scores.append(cv_score.mean())
```

```
In [138]: cv_scores
```


```
Out[138]: [0.6212624584717608,
0.6729789590254706,
0.6498338870431895,
0.6310077519379845,
0.6497231450719824,
0.6261351052048727,
0.6215946843853821,
0.6169435215946844,
0.6169435215946844,
0.6122923588039867,
0.5982281284606865,
0.5983388704318936,
0.5982281284606865,
0.5796234772978959,
0.5703211517165006,
0.5656699889258029,
0.5843853820598006,
0.5939091915836101,
0.5890365448504983,
0.56578073089701,
0.5656699889258029,
0.5611295681063122,
0.5704318936877076,
0.5703211517165006]
```

```
In [139]: neighbours[cv_scores.index(max(cv_scores))]
```

```
Out[139]: 2
```

Visualizing the K neighbors wrt CV

```
In [140]: plt.plot(neighbors, cv_scores)
plt.show()
```



```
In [141]: knn_model = KNeighborsClassifier(n_neighbors = 2)
knn_model.fit(X_train,y_train)
y_test_pred = knn_model.predict(X_test)
print("Accuracy score : " , round(accuracy_score(y_test,y_test_pred),4))
```

Accuracy score : 0.814

```
In [ ]:
```